A close-up photograph of a Maxwell GPU die mounted in a silver carrier on a black printed circuit board. The die itself is a green square with a complex grid of circuitry. The carrier has four screws at the corners. The background shows other components on the board, including an NVIDIA logo.


# Parallel Programming Using CUDA

**MAXWELL**


MOST ADVANCED GPU EVER MADE

Aman Sethi  
Debajyoti De

# Outline

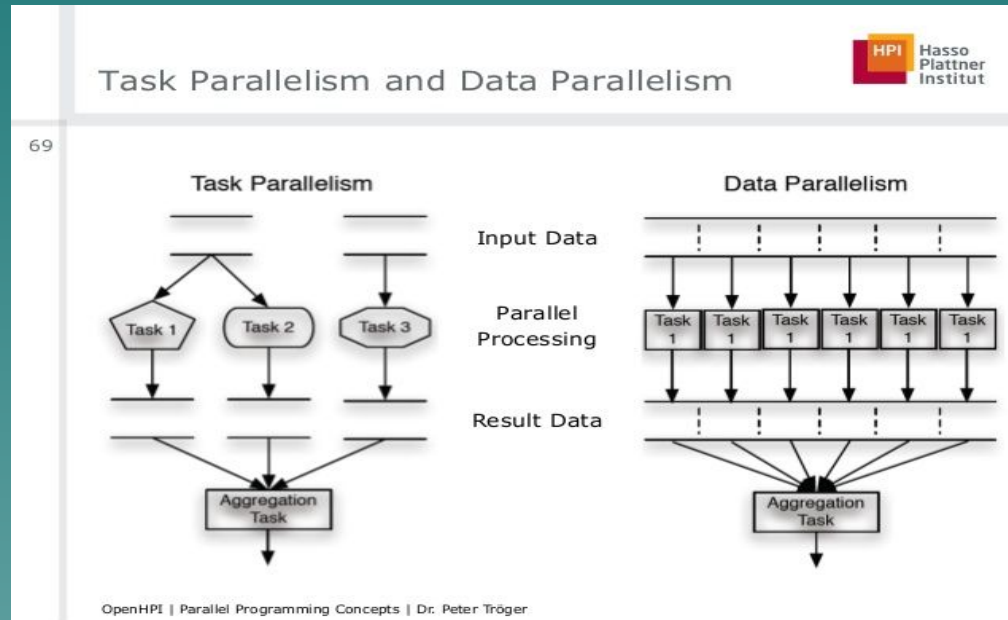
- ◆ Parallel Programming
  - ◆ GPU Programming using CUDA
    - What is CUDA?
    - How CUDA works?
    - GPU Memory Hierarchy
    - CUDA Syntax
  - ◆ Minimum in an array using CUDA
    - Approach
    - Results
  - ◆ GPU based Matrix Inversion using CUDA
    - Approach
    - Results
  - ◆ Applications
    - Hill Cipher
  - ◆ Conclusion and Further Works
- 
- A stylized, layered mountain range graphic in shades of teal and blue, located in the bottom right corner of the slide.

# Parallel Programming


- ◆ In computing, a parallel programming model is an abstraction of parallel computer architecture, with which it is convenient to express algorithms and their composition in programs.
  - ◆ A parallel program is composed of simultaneously executing processes. There are two different approaches of achieving parallelism while solving a problem.
  - ◆ They are CONTROL PARALLEL approach and DATA PARALLEL approach.
- 
- A decorative silhouette of a mountain range is located in the bottom right corner of the slide, rendered in a darker shade of the background teal color.

# Parallel Programming(continue)

- ◆ In Control Parallel approach, parallelism is achieved by applying different operations to the data set simultaneously.
- ◆ In Data Parallel approach, parallelism is achieved by applying same operation to the different data set simultaneously.



# Multithreaded Programming – A Type of Parallel Programming

- ◆ The threads model of parallel programming is one in which a single process (a single program) can spawn multiple, concurrent "threads" (sub-programs).
  - ◆ Each thread runs independently of the others, although they can all access the same shared memory space.
  - ◆ Threads can be used to implement both the data parallel and control parallel approach.
  - ◆ A challenge of using threads is the issue of collisions and race conditions, which can be addressed using synchronization.
- 

# GPU Programming Using CUDA

GPU computing is possible because today's GPU does much more than render graphics: It sizzles with a teraflop of floating point performance and crunches application tasks designed for anything from finance to medicine.

# What is CUDA??

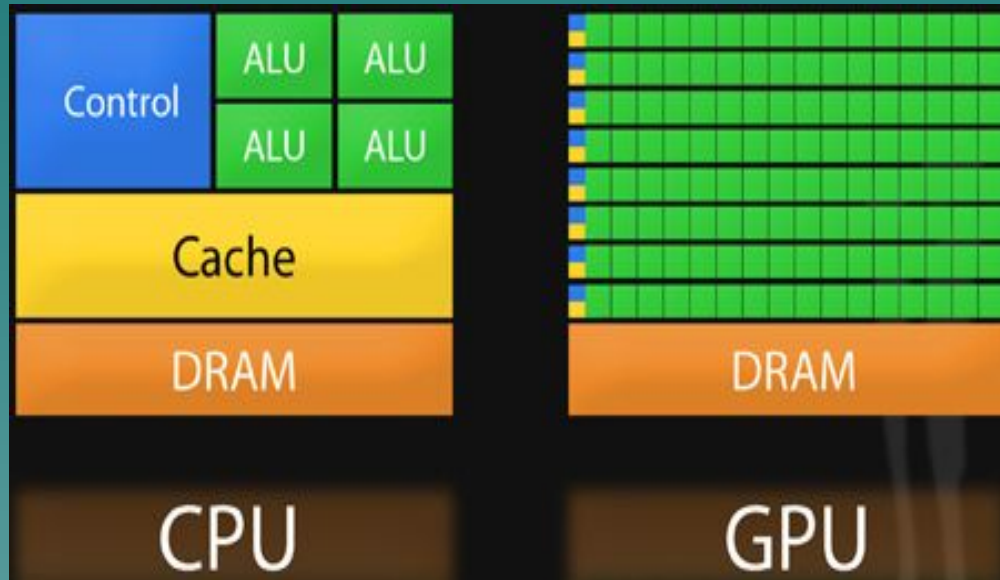
- Parallel computing architecture developed by NVIDIA
- CUDA programming interface consists of :
  - C language extensions to target portions of source code on the compute device
  - A runtime library split into:
    - Host (CPU) component
    - Device (GPU) component
    - Common component



# How CUDA works??

Data-parallel computing maps well to GPUs.

Identical operations on many data elements are performed in parallel.





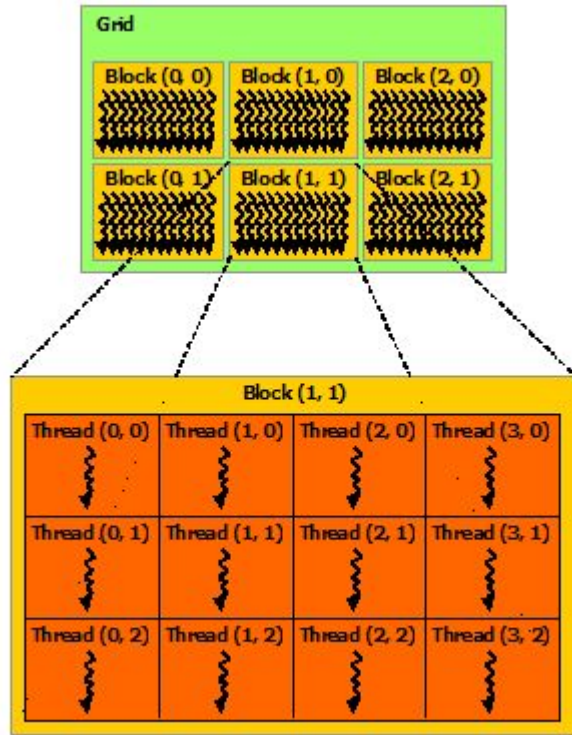
# How CUDA works??

A usual template followed to program a GPU is as below:

- Transfer the data to the GPU memory.
- Launch the GPU kernel from the host program.
- Wait till GPU finishes off the calculations.
- Transfer the results back to CPU memory.

Thus, a CUDA program consists of two parts in its code: a) **host program** which runs on CPU and b) **kernel** which does the calculations on GPU

Global, constant, texture memories



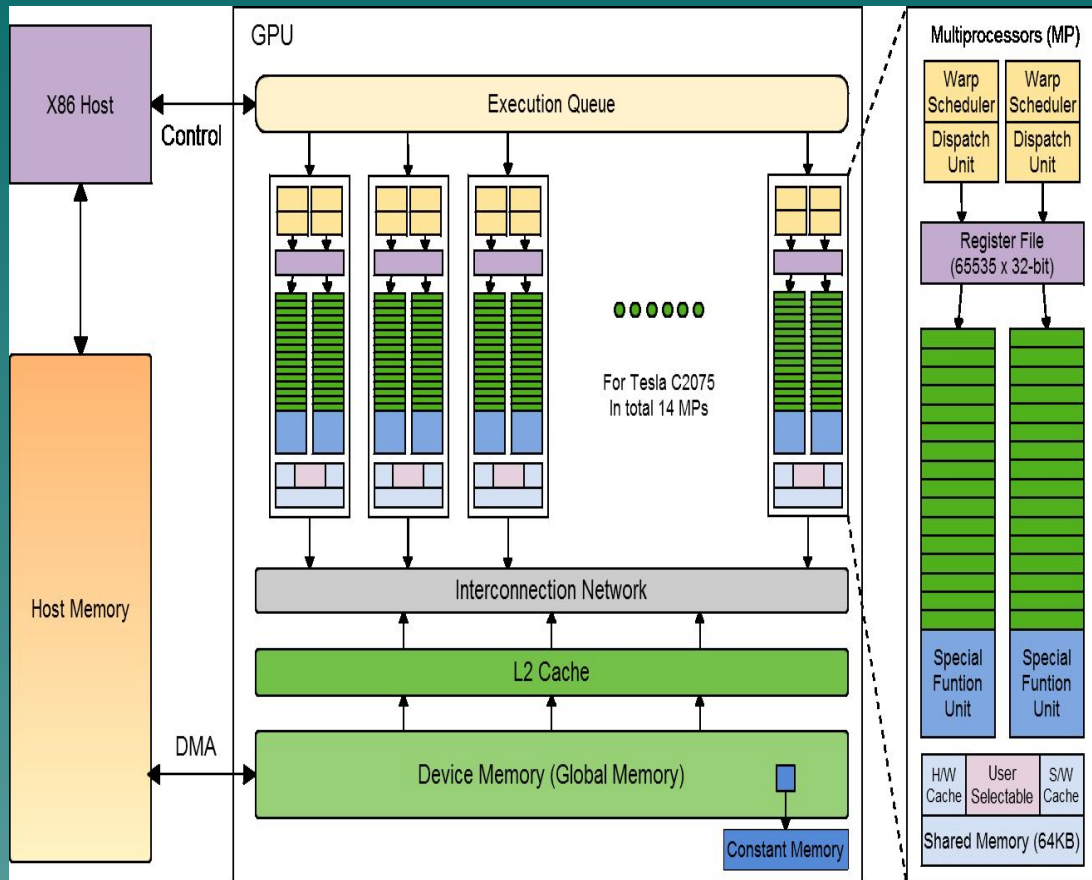
The execution of the kernels is done in threads and blocks.

A group of threads forms a thread Block.

A group of thread blocks forms a Grid.

Each thread performs same operations on a subset of data structure.

Threads within a block can synchronize using a barrier synchronization function **`__syncthreads()`**.

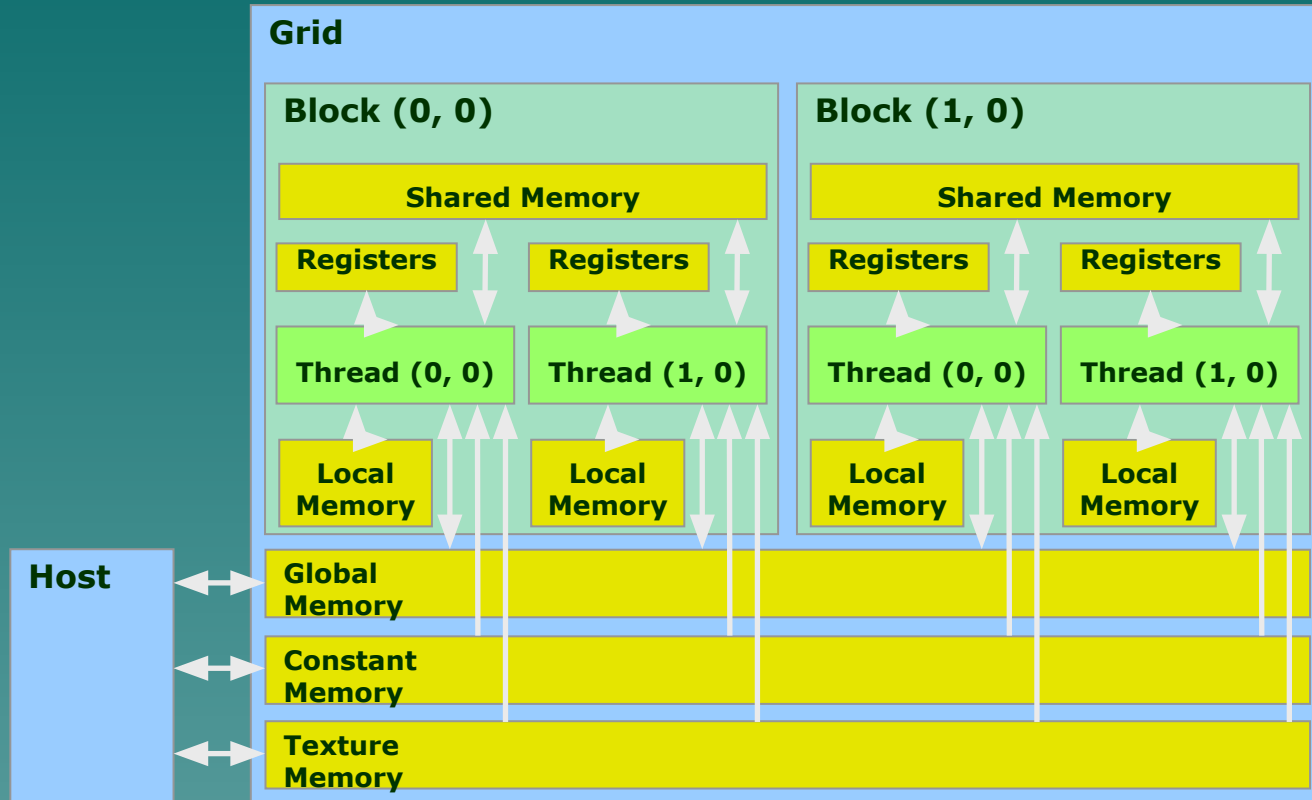


Blocks from the grid are distributed across streaming multiprocessors.

Each multiprocessor contains:

- Processor cores
- Scheduler
- Dispatcher
- Registers
- Caches
- Shared Memory

# GPU Memory Hierarchy



# GPU MEMORY MODEL SUMMARY

Memory Space	Managed By	Physical Implementation	Scope on GPU	Scope on CPU	Lifetime
Thread-Private Variables(Registers)	Compiler	On-Chip	Per Thread	Not visible	Lifetime of a thread
Thread-Private Variables(Local)	Compiler	Device Memory	Per Thread	Not visible	
Shared	Programmer	On-Chip	Block	Not visible	Block Lifetime
Global	Programmer	Device Memory	All threads	Read/Write	Application or until explicitly freed
Constant	Programmer	Device Memory	All threads	Read/Write	

# CUDA Kernels

- CUDA kernels are launched by the host using modified C function call syntax.
  - Eg: `vectorAdd <<< gridDimension , blockDim >>> (A,B,C)`
- Denoted with `__global__` function qualifier
  - Eg: `__global__ vectorAdd <<<....,.....>>> (...)`
- CUDA Kernels have few restrictions
- Kernels have read-only built-in variables:
  - `gridDim` , `blockDim` : Dimensions of grid and thread block respectively
  - `blockIdx` : Unique index of block within grid
  - `threadIdx` : Unique index of thread within thread block
- Cannot vary size of blocks or grids during kernel call

# CUDA Syntax

- Memory Management : Host code manages device memory.
  - `cudaError_t cudaMalloc(void** devPtr, size_t size)`
  - `cudaError_t cudaMemcpy(void** devPtr, int value, size_t count)`
  - `cudaError_t cudaFree(void* devPtr)`
- Data Transfer : Host code manages transfer of data to and from device.
  - `cudaError_t cudaMemcpy(void* dst, const void* src, size_t count, enum cudaMemcpyKind kind)`
- Must exercise caution when dereferencing pointers



# Minimum in an Array Using CUDA



# Approach

## On CPU(Host)

```
cudaEventCreate(&start);  
cudaEventCreate(&stop);
```

```
// Allocate memory for array a and minarr  
a = (char*)malloc(SIZE*sizeof(char));  
minarr = (char*)malloc(NO_OF_BLOCKS*sizeof(char));
```

```
// Allocate memory on the GPU for array a and minarr  
cudaMalloc(&d_a, SIZE*sizeof(char));  
cudaMalloc(&d_minarr, NO_OF_BLOCKS*sizeof(char));
```

```
/* Copy the values in array a and minarr to the memory space allocated in the GPU */  
cudaMemcpy(d_a, a, SIZE*sizeof(char), cudaMemcpyHostToDevice);  
cudaMemcpy(d_minarr, minarr, NO_OF_BLOCKS*sizeof(char), cudaMemcpyHostToDevice);
```

```
cudaEventRecord(start);
```

```
// Launching the kernel  
FindMin <<< NO_OF_BLOCKS, NO_OF_THREADS >>>(d_a,d_minarr, SIZE);  
cudaThreadSynchronize();  
cudaEventRecord(stop);
```

```
// The result array d_minarr is copied to minarr  
cudaMemcpy(minarr, d_minarr, NO_OF_BLOCKS * sizeof(char), cudaMemcpyDeviceToHost);
```

```
cudaMemcpy(d_minarr, minarr, NO_OF_BLOCKS*sizeof(char), cudaMemcpyHostToDevice);
```

```
cudaEventRecord(start);
```

```
// Launching the kernel
```

```
FindMin <<< NO_OF_BLOCKS, NO_OF_THREADS >>>(d_a,d_minarr, SIZE);
```

```
cudaThreadSynchronize();
```

```
cudaEventRecord(stop);
```

```
// The result array d_minarr is copied to minarr
```

```
cudaMemcpy(minarr, d_minarr, NO_OF_BLOCKS *sizeof(char), cudaMemcpyDeviceToHost);
```

```
min = minarr[0];
```

```
// Finding the global minimum form minarr array.
```

```
for (i = 1; i < NO_OF_BLOCKS; i++)
```

```
{
```

```
    if (minarr[i] < min)
```

```
        min = minarr[i];
```

```
}
```

```
cudaEventSynchronize(stop);
```

```
cudaEventElapsedTime(&milliseconds, start, stop);
```

```
/* Free the memory which was allocated for a and minarr in the host's memory */
```

```
free(a);
```

```
free(minarr);
```

```
/* Free the memory which was allocated for a and minarr in the GPU */
```

```
cudaFree(d_a);
```

```
cudaFree(d_minarr);
```

# On GPU(Device)

```
__global__ void FindMin(char *a,char *b, int n)
{
    // Cache array is shared by threads in a thread block.
    __shared__ char cache[NO_OF_THREADS];
    int i = threadIdx.x + blockIdx.x * blockDim.x;

    int cacheIndex = threadIdx.x;
    cache[cacheIndex] = CHAR_MAX;

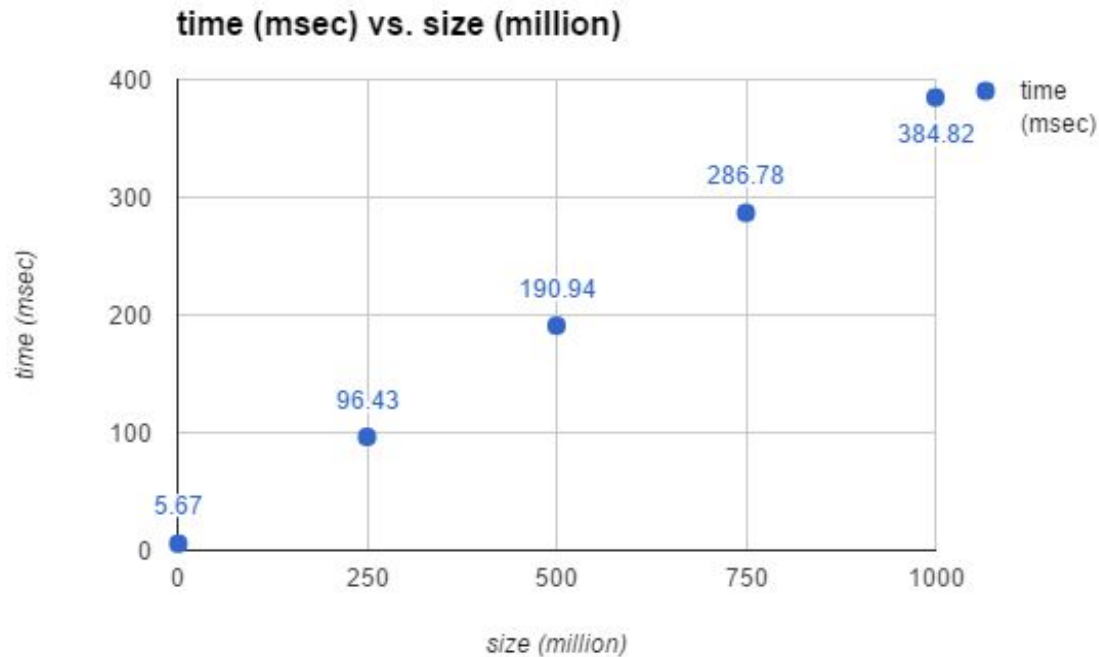
    for (int j = i; j < n;)
    {
        if (a[j] < cache[cacheIndex])
            cache[cacheIndex] = a[j];

        j += blockDim.x * gridDim.x;
    }

    /* A thread cannot proceed further unless all other threads reaches this point. */
    __syncthreads();
    /* Finding the minimum element from an array (cache) where each
    thread in a thread block has calculated its own minimum
    element */
    for (int j = 0; j < NO_OF_THREADS; j++)
    {
        if (ans > cache[j])
            ans = cache[j];
    }
    // Saving the minimum element calculated by the thread block.
    b[blockIdx.x] = ans;
}
```

# Results

The execution time was recorded for different size of the array and the following graph was obtained.



# GPU based Matrix Inversion using CUDA



# Approach

- A variant of Gaussian elimination called Gauss–Jordan elimination can be used for finding the inverse of a matrix, if it exists.
- The  $n$  by  $n$  identity matrix is augmented to the right of  $A$ , forming a  $n$  by  $2n$  block matrix  $[A \mid I]$ .
- Now through application of elementary row operations, find the reduced echelon form of this  $n$  by  $2n$  matrix.
- The matrix  $A$  is invertible if and only if the left block can be reduced to the identity matrix  $I$ ; in this case the right block of the final matrix is  $A^{-1}$



# On CPU(Host)

```
dim3 threads(NO_OF_THREADS_X, NO_OF_THREADS_Y);  
dim3 blocks(NO_OF_BLOCKS_X, NO_OF_BLOCKS_Y);
```

```
testColumnMatrix = new double[row];  
factor = new double[row];
```

```
//Allocate memory in host for matrix  
matrix = new double[(row*column)];
```

```
//Allocate memory in device for matrix  
cudaMalloc(&d_matrix, (row*column)*sizeof(double));  
cudaMalloc(&d_factor, (row)*sizeof(double));
```

```
//Initialize 2D Matrix  
initializeMatrix(matrix, row, column);  
cudaEvent_t start, stop;  
cudaEventCreate(&start);  
cudaEventCreate(&stop);  
cudaEventRecord(start);  
cudaMemcpy(d_matrix, matrix, (row*column)*sizeof(double), cudaMemcpyHostToDevice);  
for (int i = 0; i < row; i++)  
{
```

```
    // Copy ith column from matrix stored in device to host to check whether Pivot element (ith element of ith row) is zero  
    cudaMemcpy2D(testColumnMatrix, sizeof(double), &d_matrix[i], column*sizeof(double), sizeof(double), row,
```

```
cudaMemcpyDeviceToHost);
```

```
    // If pivot element is zero find a row with non-zero pivot element and interchange row elements.
```

```
    if (testColumnMatrix[i] == 0.0)
```

```
    {
```

```
        int rowNumber = findRow(testColumnMatrix, row, i);
```

```
if (rowNumber == -1)
{
    printf("Given Matrix Impossible to invert. %d\n",i);
    return 0;
}
else
{
    // Interchange rows if required
    switchRows << <numberOfBlocks, threadsPerBlock >> >(d_matrix, i, rowNumber, column);
    cudaDeviceSynchronize();
}
}
// Copy ith column from matrix stored in device to host
cudaMemcpy2D(testColumnMatrix, sizeof(double), &d_matrix[i], column*sizeof(double), sizeof(double), row,
cudaMemcpyDeviceToHost);

// Normalize Pivot ( Make pivot element 1 )
double pivotValue = testColumnMatrix[i];
normalizePivot << <numberOfBlocks, threadsPerBlock >> >(d_matrix, column, i, pivotValue);
cudaDeviceSynchronize();
// Update other rows.
for (int k = 0; k < row; k++)
    factor[k] = testColumnMatrix[k];
    cudaMalloc(&d_factor, (row)*sizeof(double));
    cudaMemcpy(d_factor, factor, row*sizeof(double), cudaMemcpyHostToDevice);
    updateOtherRows << <blocks, threads >> >(d_matrix, row, column, i, d_factor);
    cudaDeviceSynchronize();
}
cudaMemcpy(matrix, d_matrix, (row*column)*sizeof(double), cudaMemcpyDeviceToHost);
cudaEventRecord(stop);
cudaEventSynchronize(stop);
cudaEventElapsedTime(&total, start, stop);

cudaFree(d_matrix);
cudaFree(d_factor);
```

# On GPU(Device)

// The parallelised function switches the values of two rows

```
__global__ void switchRows(double *d_matrix, int i, int rowNumber, int n)
{
    int id = blockIdx.x*blockDim.x + threadIdx.x;
    if (id < n)
    {
        double temp = d_matrix[Index(i, id, n)];
        d_matrix[Index(i, id, n)] = d_matrix[Index(rowNumber, id, n)];
        d_matrix[Index(rowNumber, id, n)] = temp;
    }
}
```

// The parallelised function makes the pivot element value equal to one

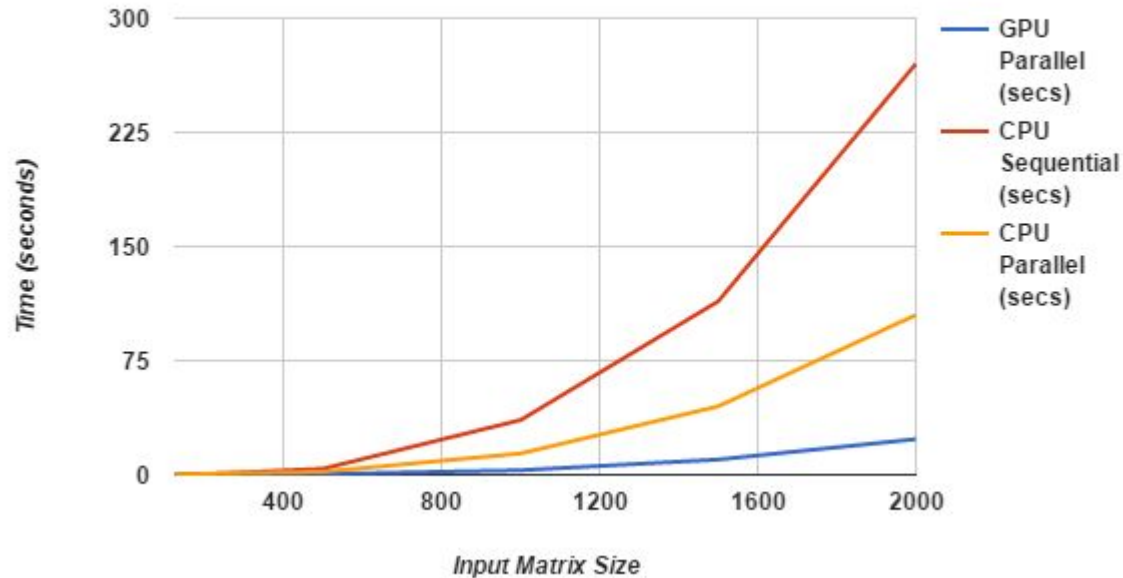
```
__global__ void normalizePivot(double *d_matrix, int n, int rowNumber, double pivotValue)
{
    int id = blockIdx.x*blockDim.x + threadIdx.x;
    if (id < n)
    {
        d_matrix[Index(rowNumber, id, n)] = (d_matrix[Index(rowNumber, id, n)] / pivotValue) ;
    }
}
```

// The parallelised function updates values of other rows so that elements above and below pivot element are zero

```
__global__ void updateOtherRows(double *d_matrix, int R, int C, int i, double *d_factor)
{
    int row = blockIdx.y*blockDim.y + threadIdx.y;
    int column = blockIdx.x*blockDim.x + threadIdx.x;
    if (row < R && row != i)
        if (column < C)
            d_matrix[Index(row, column, C)] = (d_matrix[Index(row, column, C)] - (d_factor[row] * d_matrix[Index(i, column, C)]));
}
```

# Results

- ♦ The execution time for sequential program in CPU and parallel program in CPU as well as in GPU using CUDA for different matrix size are recorded.



# Conclusion

- ◆ The execution time of matrix inversion reduces to almost half when we implement the Gauss-Jordan Elimination algorithm using multiple threads.
- ◆ The execution time of the parallel program using CUDA in GPU takes less than one-tenth of the time taken by a sequential program in a CPU when the order of the input matrix size is  $2000 \times 2000$ .

# Application

The matrix inversion is required in:

- ◆ image reconstruction process
- ◆ Hill Cipher to obtain the inverse of a key matrix during the decryption process to get back the original plaintext.

# Reference

- CUDA By Example
  - Jason Sanders
  - Edward Kandrot
- Programming GPU with CPU
  - Hasindu Gamaarachchi
- GTC Express Webinar: The CUDA Memory Model
- <http://cuda-programming.blogspot.in>



THANK YOU

