

IMPLEMENTATION OF CHANDY LAMPORT'S GLOBAL SNAPSHOT ALGORITHM USING MPI

BY

SAURA MANDAL
&
TARUN NASKAR

UNDER THE GUIDANCE OF

MR. UTPAL RAY
ASSISTANT PROFESSOR
DEPARTMENT OF INFORMATION TECHNOLOGY
JADAVPUR UNIVERSITY

OUTLINE

- ◉ INTRODUCTION
- ◉ MULTICOMPUTING AND MPI BASICS
- ◉ MONEY TRANSACTION SYSTEM
- ◉ CHANDY LAMPORT'S GLOBAL SNAPSHOT ALGORITHM
- ◉ IMPLEMENTATION OF CHANDY LAMPORT'S GLOBAL SNAPSHOT ALGORITHM USING MPI
- ◉ CONCLUSION
- ◉ REFERENCES

INTRODUCTION

- ◉ The global state of a distributed system is a collection of the local states of its components.
- ◉ Recording the global state of a distributed system is an important paradigm.
- ◉ A continuously running distributed system (Money Transaction System) is required to take global snapshots.
- ◉ The Global snapshot model is implemented using MPI in a multicomputing environment.

MULTICOMPUTING AND MPI BASICS

- ◉ Distributed memory is used in massively parallel multicomputers and provides high levels of performance.
- ◉ Multicomputers solve the Grand Challenge computational science problems.
- ◉ Multicomputers communicate by message passing.
- ◉ MPI stands for Message Passing Interface.
- ◉ MPI is a library specification for message-passing, proposed as a standard by a broadly based committee of vendors, implementers, and users.

BASIC MPI ROUTINES

- ◉ **MPI_Init()** starts the MPI runtime environment.
- ◉ **MPI_Comm_size()** gets the number of processes, N_p .
- ◉ **MPI_Comm_rank()** gets the process ID of the current process which is between **0** and $N_p - 1$, inclusive. (These last two routines are typically called right after **MPI_Init()**).
- ◉ **MPI_Send()** sends a message from the current process to another process (the destination).
- ◉ **MPI_Recv()** receives a message on the current process from another process (the source).
- ◉ **MPI_Finalize()** shuts down the MPI runtime environment.

COMPILING AND RUNNING A MPI PROGRAM

⦿ Compiling a MPI program

```
mpicc <filename> -o <objectfilename>
```

⦿ Running a MPI program

```
mpirun -np n <objectfilename>
```

n= The number of processors.



Run an object file on multiple machines using a command:

```
mpirun -np n -pernode --hostfile my_host <objectfilename>
```

pernode ensures a single process executes in each node.

hostfile contains the IP addresses of the nodes. The first IP address of the hostfile is the master processor's IP address.

MONEY TRANSACTION SYSTEM

- A money transaction system is a system of a number of nodes (N), implemented to transfer a random amount of money from any node to any other node in a distributed environment.
- Checks data consistency in global snapshot algorithm.

CHANDY LAMPORT'S GLOBAL SNAPSHOT ALGORITHM

Algorithm:

Marker Sending Rule for process i :

begin

- (i) Process i records its state.
- (ii) For each outgoing channel C on which a marker has not been sent, i sends a marker along C before i sends further messages along C .

end

Marker Receiving Rule for process j :

On receiving a marker along channel C :

If j has not recorded its state **then**

begin

- (i) Record the state of C as the empty set.
- (ii) Follow the '**Marker Sending Rule**'.

end

else

Record the state of C as the set of messages received along C after j 's state was recorded and before j received the marker along C .

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

Assumptions: **Recording State** and **Sending Marker** are **atomic** operations.

Each processor is initialized with \$10,000.

The below code snippet does the money transaction system.

```
for (int i=0;i<size;i++) {  
    if(rank==i)  
        continue;  
    int d=rand()%10+1;  
    MPI_Send(&d,1,MPI_INT,i,0,MPI_COMM_WORLD);  
    data-=d;  
    g_node.channel_record[i]=data;  
    g_node.channel_trans[i]=d;  
    MPI_Recv(&d,1,MPI_INT,i,0,MPI_COMM_WORLD,&status);  
    data+=d;  
}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
//initialize the channel buffer, state and status of channel
for(int i=0;i<size;i++) {
    g_node.channel_status[i]=RED;
    g_node.channel_record[i]=null;
}
g_node.state=data;
srand(time(NULL));
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
while (true) {  
    perform_money_system_transaction();    //running indefinitely  
    scanf("%c %d",&MARKER,&initiator);  
    if(initiator > size-1) {  
        printf("Enter a valid node!\n");  
        MPI_Abort(MPI_COMM_WORLD,1);  
    }  
    //SENDING MARKER  
    for(int i=0;i<size;i++) {  
        if(rank==i)  
            continue;  
        else if(rank==initiator) {  
            //start recording on incoming channels and make incoming channels wideopen  
            for(int j=1;j<size;j++)  
                g_node.channel_status[j]=RED;  
        }  
    }
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
//send markers to other processors
ierr=MPI_Isend(&MARKER,1,MPI_CHAR,i,0,MPI_COMM_WORLD,&send_request);
    g_node.state=data;    //record own state
    ierr=MPI_Wait(&send_request,&status);
}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
//RECEIVING MARKER
for(int i=0;i<size;i++) {
    if(rank==i)
        continue;

    //received marker for the first time

    ierr=MPI_Irecv(&MARKER,1,MPI_CHAR,i,0,MPI_COMM_WORLD,&recv_request);
    ierr=MPI_Wait(&recv_request,&status);
    //if marker is received for the first time, record own state
    if(g_node.channel_status[i]==RED) {
        g_node.state=data;          //record state
        g_node.channel_record[rank]=0; //initiate NULL
        g_node.channel_status[i]=GREEN; //close the channel
    }
}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
else { //if received earlier
```

```
    if(rank==initiator) {  
        if(g_node.channel_status[i]==GREEN)  
            ;  
        //check if all the incoming channels are recorded GREEN  
        for(int j=0;j<size;j++) {  
            if(j==initiator)  
                continue;  
            if(g_node.channel_status[j]==RED) {  
                g_node.allgreen=0;  
                break;  
            }  
        }  
    }  
}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
//if all are green we are done
    if(g_node.allgreen)
        break;
    else {
        g_node.allgreen=true;
        break;
    }
}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
else {  
    //for non-initiator processes check if all are green  
    for(int j=0;j<size;j++) {  
        if(rank==j)  
            continue;  
        if(g_node.channel_status[j]==RED) {  
            g_node.allgreen=0;  
            break;  
        }  
    }  
    if(!g_node.allgreen)  
        continue;  
    else {  
        g_node.allgreen=true;  
        break;  
    }  
}}}
```


IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
for(int i=0;i<size;i++)
    g_node.state-=(g_node.channel_trans[i]);
sum+=g_node.state;

array[0]=g_node.state;
for(int i=1;i<(size+2-1);i++)
    array[i]=g_node.channel_trans[i-1];
array[size+2-1]=sum;
MPI_Send(&array,size+2,MPI_INT,0,0,MPI_COMM_WORLD);

}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
if(rank==0) {  
  
    for(int i=0;i<size;i++) {  
        sum+=g_node.channel_trans[i];  
    }  
  
    for(int i=0;i<size;i++)  
        g_node.state-=(g_node.channel_trans[i]);  
    sum+=g_node.state;  
}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
int f=0;
f+=sum;
int i;

for(i=1;i<size;i++) {
    MPI_Recv(&array,size+2,MPI_INT,i,0,MPI_COMM_WORLD,&status);
    printf("\t%d",i);
    for(int j=1;j<(size+2);j++)
    {
        //info[i][j-1]=tarray[j-1];
        printf("\t%d",array[j-1]);
    }
}
```

IMPLEMENTATION OF CHANDY LAMPORT'S ALGORITHM USING MPI

```
f+=array[size+2-1];
    if(i==size-1){
        printf("\n\n\n\n INITIATOR PROCESSOR: %d",initiator);
        printf("\n\n TOTAL SUM: %d",f);
        totaltime = ((double) (final_time - initial_time));
        printf("\n\n TIME TAKEN : %f sec\n\n",totaltime);
    }
    else
        printf("\n");
}
}
```

OBSERVATIONS

m 0

CHANDY	LAMPORTS	GLOBAL		SNAPSHOT	
PROCESSOR	STATE	P0	P1	P2	P3
0	9975	0	9	10	3
1	9980	4	0	6	3
2	10010	1	3	0	5
3	9974	7	2	8	0

INITIATOR PROCESSOR: 0

TOTAL SUM: 40000

TIME TAKEN : 0.000082 sec

The final sum is always consistent.

CONCLUSION

- ◉ The basic idea is to use two colors viz. **RED** and **GREEN** which indicate whether a process has already taken its local snapshot and whether a message was sent before or after the local snapshot of a process.
- ◉ Messages which would make a snapshot inconsistent can easily be recognized and avoided, and messages which are in transit can be caught by the receiving process.
- ◉ The total amount of money was always consistent.

FURTHER WORK

- ◉ There are several variants of the Chandy Lamport's snapshot algorithm. Like **Spezialetti** , **Venkatesan's** and **Kearns** algorithm.
- ◉ More distributed systems like the money transaction system can be developed to check these algorithms.

References

- ◉ ONLINE RESOURCE: MPI-1 standard,
<http://www.mpi-forum.org/docs/mpi-1.1-html/mpi-report.html>
- ◉ Chandy K M and Lamport L “Distributed snapshots: determining global states of distributed systems” **ACM Trans. Computer Systems** 3 1985 ,pp 63-75
- ◉ Acharya A and Badrinath B R “Recording distributed snapshots based on causal order of message delivery” **Information Processing Letters** 1992, pp 317-21
- ◉ Spezialetti M and Kearns P “Efficient distributed snapshots” **Proc. 6th Int. Conf. on Distributed Computing Systems**, 1986, pp 382

THANK YOU